

From Kernel to Karma: Integrating Biophysical Models for the Safe Deployment and Autonomous Security of Cyboquatic Air-Globes

This report presents a comprehensive deep research analysis for the development of new mathematical models to support the planning, safety, and security of cyboquatic air-globes. These conceptual devices purify air using tightly contained fire-and-water mechanisms. The primary directive of this research is to create models that prioritize integration with the established CEIM (Cyber-Ecological Integrity Model) and CPVM (Cyber-Physical Viability Model) kernels, extending their functionality rather than replacing them. This approach ensures continuity with existing governance systems and maintains comparability with data from other environmental nodes, such as those managing PFBS and E. coli reductions. The framework is designed to serve two critical functions: first, to provide quantitative guidance for physical deployment and scaling, particularly for securing acceptance of a pilot program in Phoenix; and second, to enable robust, real-time operational security and autonomy for individual air-globe units. The mathematical forms are designed to be universally applicable, allowing different regions to substitute their own climatic, regulatory, and grid data while reusing the same core equations, thus ensuring broad generalizability. This report details the theoretical underpinnings, mathematical formulations, practical implications, and implementation strategies for three key models: the Eco-Safety Phase Diagram (ESPD), the Multonry Sensor-Trust Displacement scalar (D_t), and the Karma-Tolerance Security Field (KTSF). It further outlines the data schema required for governance and provides strategic recommendations for model development.

The Eco-Safety Phase Diagram (ESPD): A Quantitative Tool for Deployment Planning

The Eco-Safety Phase Diagram (ESPD) is a novel mathematical model designed to address the immediate need for quantitative, high-level guidance in the planning and

deployment of cyboquatic air-globe arrays . Its primary function is to translate complex biophysical performance metrics into a simple, intuitive decision space for planners and regulators, directly facilitating the goal of gaining acceptance for Phoenix pilots . The ESPD operates by treating each potential air-globe installation as a point within a two-dimensional phase plane, where the axes represent "eco-benefit" (B) and "risk-of-harm" (R) . By partitioning this plane into distinct zones—"Deployable," "Pilot-only," and "Forbidden"—the model provides a clear, quantitative rule set for siting decisions, scaling limits, and material selection criteria . This approach moves beyond qualitative assessments, grounding deployment strategy in a formalized, auditable mathematical framework that is consistent with the overarching CEIM/CPVM architecture .

The core of the ESPD model lies in its mathematical definitions for the eco-benefit (B) and risk (R) axes, which are explicitly designed to be extensions of existing, proven calculations within the CEIM system . The eco-benefit axis, B, quantifies the net positive impact of an air-globe's operation. It is defined by the formula:

$$B = \frac{M_{\text{captured}} - M_{\text{embodied}} - M_{\text{power}}}{M_{\text{ref}}}$$

In this equation, each mass term (M) is calculated using the standard CEIM mass-balance kernel, $M = (C_{\text{in}} - C_{\text{out}})Qt$, which has been validated for pollutants like PFBS and E. coli in Arizona water nodes . This direct linkage is a cornerstone of the integration philosophy. M_{captured} represents the mass of target pollutants (e.g., CO₂, PM) removed from the environment. M_{embodied} accounts for the mass of materials used in the construction of the air-globe itself, including sorbents and structural components. M_{power} represents the mass of carbon dioxide emissions associated with the energy consumed during operation. Finally, M_{ref} is a normalization factor, typically a reference mass, which renders the benefit score dimensionless and comparable across different scales and node types . The use of the established CEIM kernel for all mass terms guarantees dimensional consistency and allows the air-globe's performance to be integrated directly into the broader ecological accounting system, enabling comparisons with other nodes like PFBS/E. coli reduction systems . For instance, the physical plausibility of capturing 200 tons of CO₂ per year at building-scale flows is supported by the known mass conversion factor for CO₂ at standard temperature and pressure, approximately 1.9×10^{-6} kg per ppm per m³, aligning with the corrected Air-Globe kernel .

The risk-of-harm axis, R, is designed to quantify the potential negative impacts of an air-globe's operation. It is formulated as a convex fusion of multiple harm sources, represented by the equation:

$$R = wV(1-V) + w_m R_{\text{materials}} + w_n R_{\text{noise}} + w_s R_{\text{siting}}$$

Here, V is the viability scalar from the CPVM, a value in the range [0,1] that indicates how closely the system's state trajectory adheres to a safe operating set . A value of $V=1$ signifies perfect adherence to the safe set, while values closer to 0 indicate deviation and potential instability. The term $(1-V)$ therefore serves as a direct measure of the dynamic safety risk, making it a crucial input for the ESPD. The other terms in the equation account for static or external risks. $R_{\text{materials}}$ is a normalized index representing the ecotoxicological risk posed by the materials used in the device, particularly the sorbent media and any wash liquids. This component is grounded in empirical data from standardized tests. R_{noise} and R_{siting} represent the acoustic footprint and the burden on surrounding infrastructure or communities, respectively. The weights (wV, w_m, w_n, w_s) are non-negative coefficients that sum to one, allowing planners to adjust the relative importance of each risk type based on the specific context of the deployment site . The inclusion of the CPVM viability scalar V creates a powerful feedback loop, ensuring that operational safety, governed by the CPVM, directly influences high-level deployment decisions. This prevents deployments that might appear beneficial on paper but are associated with high operational risk. The overall risk score R is also a dimensionless quantity suitable for plotting on the same scale as the benefit score B .

The true power of the ESPD emerges when the (B, R) plane is partitioned into distinct zones. This partitioning is not arbitrary; it is informed by the calibration of the models against Phoenix-specific data and standards. For example, the constraint that the grid's carbon intensity must be below 50 g CO₂/kWh to ensure a net-positive removal balance for DAC-style nodes is a hard-coded deployment criterion derived from CEIM analysis of the Phoenix grid . This threshold would define a boundary within the phase diagram. Similarly, the ecotoxicological profiles of materials, validated against OECD guidelines, would inform the acceptable range for the $R_{\text{materials}}$ term . The resulting zones provide clear directives:

- **Deployable Zone:** Areas with high benefit (B) and low risk (R). Nodes falling here are deemed ready for full-scale deployment.
- **Pilot-only Zone:** Areas with moderate benefit and/or risk. This zone would require additional monitoring, specific operational protocols, or limited-scale testing before wider rollout.
- **Forbidden Zone:** Areas with low or negative benefit (B) or unacceptably high risk (R). Nodes falling here would be prohibited from deployment.

This zoning directly addresses the user's requirement for guidance on siting and scaling arrays for the Phoenix pilots . The coordinates of each node, (B_{raw} , R_{raw}), can be stored in a dedicated `qpudatashard`, providing auditable evidence for governance bodies like EcoNet and forming the basis for automated decision support systems . The model's utility is further enhanced by its ability to handle trade-offs. A planner could analyze the diagram to see if a higher-risk material ($R_{\text{materials}}$) could be justified if it leads to a significantly higher capture rate (M_{captured}), thereby increasing the B score. The weights in the risk equation allow for explicit modeling of such cost-benefit analyses. The empirical calibration points from existing Arizona nodes, which show ecoimpact scores in the 0.8–0.9 range, provide a realistic baseline for tuning the normalization factors and weightings in the model . In essence, the ESPD translates the complex interplay of physics, chemistry, biology, and engineering inherent in air-globe operation into a simple, yet powerful, graphical tool for strategic planning and governance.

Multonry Sensor-Trust Displacement (D_t): A Dynamic Guardrail for Operational Integrity

While the Eco-Safety Phase Diagram (ESPD) provides a strategic overview for deployment planning, the Multonry Sensor-Trust Displacement scalar, denoted D_t , operates at the tactical level, ensuring the integrity of real-time operations and forming a critical defense against sensor-based attacks . It is a time-windowed trust scalar that dynamically modulates the credibility of a node's reported performance metrics . Its fundamental purpose is to prevent a single faulty, mis-calibrated, or malicious sensor from inflating an air-globe's claimed eco-impact, thereby undermining the entire governance structure built upon the CEIM/CPVM kernels . By continuously assessing sensor health and cross-validating measurements against the system's underlying physical laws, D_t acts as an automatic audit trail and a dynamic guardrail for both data reporting and security enforcement.

The mathematical formulation of D_t is designed to be a comprehensive proxy for sensor reliability over a recent time window $[t-\tau, t]$. The core equation is expressed as:

$$D_t = 1 - \alpha \Delta \text{drift} - \beta \Delta \text{var} - \gamma \Delta \text{resid} - \delta N_{\text{violations}}$$

where the result is clamped to the interval $[0,1]$. Each term in this equation quantifies a different mode of sensor failure or anomaly. Δdrift measures the systematic deviation of a sensor's output from its baseline mean value. Δvar captures a blow-up in the sensor's

variance, indicating instability or noise. Δ_{resid} is derived from cross-sensor residuals, comparing readings from redundant sensors measuring the same physical quantity to identify discrepancies that violate the CEIM mass-balance kernel . Finally, $N_{\text{violations}}$ counts the number of times CPVM safety flags were triggered, which could be correlated with anomalous sensor inputs . The coefficients $(\alpha, \beta, \gamma, \delta)$ are tunable parameters that weight the impact of each diagnostic, allowing the system to prioritize certain failure modes over others. This multi-faceted approach is more robust than relying on a single metric and aligns with principles from attack-resistant trust models in wireless sensor networks (WSNs) [70](#) [71](#) . Research into dynamic calibration of low-cost sensors further supports the need for continuous, adaptive assessment of sensor health [54](#) .

Once computed, the D_t scalar is not merely an indicator of poor quality; it is an active control parameter that gates the perceived performance of the air-globe. All public-facing scores are adjusted by multiplying them with D_t . Specifically, the adjusted eco-benefit and karma scores are calculated as:

$$B_{\{\text{adj}\}} = B \cdot D_t, \quad \text{Karma} < \text{em_text_raw}="\\text{raw}">{} \{ \text{adj}\}$$

This mechanism ensures that a sensor problem can only reduce, never inflate, the credited impact of a node . If a CO₂ sensor begins to drift high, causing an overestimation of M_{captured} , the corresponding drop in D_t will automatically down-weight the reported benefit and karma, reflecting a more accurate picture of the node's actual contribution . This embeds a powerful data integrity feature directly into the governance model. Any attempt to game the system by tampering with sensor readings would be immediately penalized by a lower D_t , making such actions counterproductive.}

From a security perspective, a low D_t value is a strong signal of a compromised or failing node. The security systems can treat nodes with low trust as needing tighter controls, effectively quarantining them from the network until the issue is resolved . This makes the overall cyboquatic ecosystem resilient to targeted attacks aimed at corrupting data streams. The concept of using trust estimation as a basis for security-aware sensor fusion is well-established, often modeled using frameworks like hidden Markov models [69](#) . The D_t scalar provides a concrete, computationally efficient implementation of this principle for the air-globe context. Furthermore, the multitory nature of the model, considering multiple diagnostics simultaneously, helps to distinguish between benign anomalies (e.g., a temporary spike in noise) and systemic failures (e.g., sustained drift), reducing false positives in the security response [20](#) . The model is already defined and integrated into merged air-water shards, demonstrating its readiness for implementation . By linking the physical reality of sensor measurements to the abstract governance metrics of benefit and

karma, the Multonry Sensor-Trust Displacement scalar closes the loop between real-time operation and long-term accountability, ensuring that the system's actions are always grounded in reliable data.

Karma-Tolerance Security Fields (KTSF): Adaptive Cyber-Physical Defense for Purification Nodes

The Karma-Tolerance Security Field (KTSF) extends the abstract concept of karmatolerance, originally developed for human identities, to the machines themselves, creating a sophisticated and adaptive security posture for the cyboquatic network . This model moves beyond static access controls and one-size-fits-all security policies, instead proposing a dynamic defense mechanism where the security treatment of a node is determined by its measured behavior and reputation within the ecosystem. Each purification machine (node) is assigned a "Karma triple" comprising its eco-impact, its contribution to the network, and its security-trustworthiness, which are then fused into a single Karma score, K_i . This score, in turn, defines a "tolerance radius," T_i , which sets the upper limit on how harshly security responses may be applied to that specific node . This approach mathematically incentivizes good behavior, protects critical assets from trivial disablement, and enables rapid quarantine of compromised or anomalous nodes.

The foundation of the KTSF model is the Karma score, K_i , which is calculated using a convex fusion of its constituent parts, mirroring the existing model for human identities. The formula is given by:

$$K_i = w_E E_i + w_C C_i + w_S S_i$$

where E_i is the normalized eco-impact score, C_i is the contribution score, and S_i is the security-trust score for node i . The weights (w_E, w_C, w_S) sum to one, allowing for flexible prioritization of different aspects of a node's performance . The use of a normalized, dimensionless score ensures mathematical consistency across the entire governance system, whether applied to a human participant or a purification machine . This fusion of benefits and trust into a single metric is a core tenet of the "eco-grammar" governing the cyboquatic ecosystem . The normalized node impact metric, K_n , used in CEIM, is already known to be stable and additive, making it a suitable building block for this higher-level construct . Empirical data from existing Arizona nodes, which exhibit ecoimpactscores in the 0.8–0.9 range, can be used to empirically calibrate the scoring ranges for E_i and K_i .

The critical innovation of the KTSF model is the mapping of this Karma score to a tolerance radius, $T_i = f(K_i, B_i, R_i)$. This function determines the aggressiveness of the security response that can be directed at the node. Nodes with a high Karma score (K_i close to 1), which typically correspond to high benefit (B_i) and low risk (R_i), are granted a large tolerance radius. This means they are considered critical assets, and security actions against them—such as firmware locks, rate-limiting, or remote shutdowns—are strictly clamped and difficult to execute. This protects the network's most valuable purification resources from denial-of-service-like attacks or accidental misconfiguration. Conversely, nodes with a low Karma score, perhaps due to noisy operation, low efficiency, or a history of security violations, are assigned a small tolerance radius. For these peripheral or potentially compromised assets, security systems have greater freedom to apply aggressive quarantining measures to prevent them from becoming a vector for wider network compromise.

This adaptive security paradigm aligns with formal approaches to modeling CPS security, where the adversary's capabilities and the system's response are governed by a threat model [30](#) [31](#). The KTSF provides a concrete, computable instantiation of such a model. Reputation-weighted Intrusion Detection Systems (IDS) and STRAC-style models have empirically demonstrated that incorporating reputation into security decisions can reduce false positives and improve response accuracy, lending support to the core premise of the KTSF. The security response cap, which could be the severity of a lock, the duration of a quarantine, or the authority to perform a remote reset, is clamped by T_i . This ensures that the security posture is not uniform but is instead tailored to the perceived value and trustworthiness of each individual asset. For example, a high-karma air-globe contributing significantly to Phoenix's air quality would require a significant, auditable justification to be taken offline, whereas a low-karma experimental sensor node could be instantly isolated with minimal procedural overhead. This creates a powerful incentive structure: operators are motivated to maintain high-quality sensors (D_t), operate efficiently (high B), and adhere to safety protocols (low R) to build up their node's Karma score and, consequently, its resilience to external threats. The KTSF model thus transforms security from a reactive, perimeter-based concept into a proactive, reputation-aware capability embedded directly within the operational logic of the machines themselves.

A Generalizable Modeling Architecture Anchored in Proven Standards

A central tenet of the research goal is the creation of a modeling architecture that is both highly specialized for the Phoenix pilot and universally generalizable to other regions . This dual requirement is met by designing the mathematical forms of the new models as universal templates, while treating region-specific parameters as interchangeable inputs. This architectural choice preserves the stability and interoperability of the core CEIM and CPVM kernels while allowing the system to adapt to diverse ecological, climatic, and regulatory environments without requiring changes to the underlying code or formulas . The strategy is to anchor the models in rigorously verifiable scientific standards, ensuring that the "generalizable" aspect is not just a theoretical claim but a practice grounded in empirical evidence.

The foundation of this generalizable architecture is the preservation of the CEIM mass-balance kernel, $M=(C_{in}-C_{out})Qt$, and the CPVM viability scalar, $V\in[0,1]$, as the backbone of the new models . These elements are treated as constants. The new models —the ESPD, D_t , and KTSF—all take their outputs as inputs or are constructed directly from them. For instance, the eco-benefit (B) in the ESPD is derived from masses calculated with the CEIM kernel, and the risk metric (R) incorporates the CPVM scalar V . Because these foundational components are already designed to accept region-specific thresholds (e.g., EPA, EU, WHO standards; local NOEC/EC50 values), the new models inherit this flexibility . When deploying in a new region, the operators would not alter the equations for B or R; instead, they would substitute their local data for the parameters within those equations .

The calibration process exemplifies this generalizable design. For the initial Phoenix pilot, the models are calibrated using a specific set of Phoenix-centric facts . This includes the grid carbon intensity ceiling of ≈ 50 g CO₂/kWh, which is a hard constraint encoded in the ESPD to guarantee net-positive capture . It also includes the specific desert climate (e.g., high temperatures, alkaline water composition), documented airshed conditions, and the target annual capture capacity of 200 tons per school-sized node . These parameters are not hardcoded into the mathematics but are bundled together as a verifiable parameter set, akin to a "Phoenix evidence hex-bundle" . This bundle represents the specific configuration required for successful deployment in that corridor.

When moving to a different region, such as an urban area in Iceland or a basin in another country, the same mathematical forms for the ESPD, D_t , and KTSF are reused . The operators in the new region would procure their own "evidence hex-bundle" containing

their local equivalent data: their national grid's average carbon intensity and renewable penetration windows, local climate and airshed data, and relevant regulatory discharge benchmarks for wash liquors . The software remains unchanged; only the input data is swapped. This approach dramatically reduces the barrier to entry for new deployments and ensures consistency across the global cyboquatic network. The governance math stays stable, allowing for meaningful comparisons of performance and risk between a Phoenix air-globe and a Gila Estrella water node, for example, because both are ultimately measured against the same universal yardstick provided by CEIM and CPVM .

This entire architecture is anchored in internationally recognized scientific standards, which provides the necessary rigor and objectivity to support its generalization. The requirement for internal sorbent and material biodegradability is validated against ISO 14851, which specifies a method for determining ultimate aerobic biodegradability by measuring oxygen demand in a closed respirometer [1](#) [8](#) . The ecotoxicological risk of leachates is assessed using OECD Test Guidelines 201 (freshwater alga and cyanobacteria growth inhibition) and 202 (terrestrial plants), which provide standardized procedures for evaluating the effects of substances on aquatic and terrestrial life [6](#) [40](#) . The plan to adapt these methods to a Phoenix water matrix (e.g., pH 8.4, TDS 500–1200 mg/L) demonstrates a commitment to maintaining scientific validity even when applying the standards to a novel context . Other relevant standards include ISO 14855 for composting biodegradation and OECD 301 series tests, which offer a multi-parameter approach to validation through CO₂ evolution, O₂ consumption, and DOC measurement [3](#) . By tying the abstract model variables (like *R*_{materials}) directly to these empirical, lab-verifiable test results, the entire modeling framework gains a solid evidentiary foundation that transcends any single location, supporting its claim to be both locally adaptable and globally consistent.

Data Schema and Implementation Blueprint for Governance and Control

To translate the theoretical models of ESPD, *D_t*, and KTSF into a functional, auditable system, a standardized data schema is essential. The qpudatashard file format, specifically structured as a CSV, serves as the canonical ledger for recording the state and performance of every cyboquatic node . This shared data structure ensures that all components of the system—from deployment planners to real-time security controllers—ingest and interpret information consistently. The blueprint for this schema, derived from

the research goal and conversation history, integrates the outputs of the new mathematical models directly into the existing particle-based architecture, creating a unified source of truth for governance and control.

The proposed **PlanningSafetySecurityAirWater2026v1.csv** shard is designed to be a comprehensive record of a node's performance and status over a specific time window . It fuses the key metrics from the three new models into a single planning-grade structure, providing all necessary inputs for both high-level siting decisions and low-level security orchestration . The table below outlines the required fields for this shard, along with their descriptions and relevance to the core research goal.

Field Name	Description	Relevance to Research Goal
nodeid	Unique identifier for the node, e.g., 'AG-FW-PHX-01'.	Ensures all data is correctly attributed to a specific physical asset for traceability and governance.
medium	The medium being processed, e.g., 'air' or 'water'.	Allows the unified kernel to handle different purification contexts (air vs. water) within the same logical framework.
region	Geographic region of deployment, e.g., 'Phoenix-AZ'.	Enables regional parameter substitution and aggregation of performance data.
twindowstart / twindowend	Start and end timestamps for the reporting period.	Provides temporal granularity for tracking performance and detecting anomalies over time.
B_raw	Raw, unadjusted eco-benefit score, calculated via the ESPD model.	The core performance metric for planning, showing the theoretical benefit before trust adjustments.
R_raw	Raw, unadjusted risk-of-harm score, calculated via the ESPD model.	The core risk metric for planning, showing the theoretical risk before trust adjustments.
Dt	Time-windowed Multony Sensor-Trust Displacement scalar.	The dynamic trust score that down-weights all public claims; critical for security and data integrity.
K_i	Karma score for the node, a fusion of eco-impact, contribution, and security-trust.	The reputation metric used by the KTSF model to modulate security responses.
T_i	Karma-Tolerance Security Field radius for the node.	The security response cap, determining how aggressively the node can be controlled.
B_adj	Adjusted eco-benefit score, calculated as $B_{adj} = B_{raw} \cdot Dt$.	The final, trustworthy performance score that is publicly reported and contributes to the node's karma.
security_response_cap	The designated security action that is permissible for this node.	An operational output that guides real-time security decisions based on the KTSF model.
evidence_hex	A Merkle-linked hexadecimal string representing the evidence package for this shard.	Provides cryptographic proof of the data's origin, integrity, and the evidence used for its calculation.

This schema provides a complete blueprint for implementation. The **B_raw** and **R_raw** fields are populated by the ESPD model, using the CEIM and CPVM kernels as their primary computational engine . The **Dt** field is generated by a parallel module running

the Multonry Sensor-Trust Displacement algorithm, ingesting raw sensor telemetry and publishing the trust score . The K_i and T_i fields are calculated based on the node's historical performance, aggregating past B_{adj} and trust scores to build a persistent reputation . The `security_response_cap` is then determined by querying the KTSF model with the current T_i value .

For the Phoenix pilot, this schema is directly aligned with the existing AirGlobe and CAIN shards, requiring no schema change to implement the new models . The fields are machine-readable and ready for direct ingestion into the C++/ALN stack, ensuring seamless integration with the existing technology . The inclusion of the `evidence_hex` field is a critical feature, as it anchors each data point to a verifiable bundle of evidence, including the design recipe, compliance certificates (ISO 14851, OECD 201), and calibration data . This creates an immutable audit trail, which is essential for governance by EcoNet and for building public trust. The implementation of this blueprint would involve developing C++ libraries for each of the new models, which would read from and write to the `qudatashards` directory, ensuring that the entire system operates on a common, trusted dataset. This structured approach to data management is fundamental to realizing the goals of safe, secure, and effective deployment of cyboquatic air-globes.

Synthesis and Strategic Recommendations for Model Development

This research report has detailed the development of a suite of new mathematical models —namely the Eco-Safety Phase Diagram (ESPD), Multonry Sensor-Trust Displacement (D_t), and Karma-Tolerance Security Fields (KTSF)—designed to extend the capabilities of the existing CEIM and CPVM frameworks for cyboquatic air-globes. The overarching architectural principle is one of integration over replacement, preserving the stability of proven kernels while layering new biophysical dimensions to enhance planning, safety, and security . The proposed models successfully fulfill the dual-purpose requirement: the ESPD provides a quantitative tool for strategic deployment planning, directly addressing the near-term need for guidance in the Phoenix pilot program, while the D_t and KTSF models provide a robust, adaptive layer of real-time operational security and autonomy . Furthermore, the design is fundamentally generalizable, with universal mathematical forms that can be calibrated for any region by substituting local data, a principle anchored in internationally recognized scientific standards like ISO 14851 and OECD 201 .

The synthesis of these models reveals a cohesive and powerful framework. The ESPD acts as the strategic planner, defining the "rules of engagement" for deployment based on a trade-off between benefit and risk. The D_t scalar serves as the real-time integrity monitor, ensuring that the data flowing from each node is trustworthy and cannot be easily manipulated. Finally, the KTSF model functions as the adaptive security enforcer, tailoring defensive actions to the specific reputation and value of each node. Together, they create a multi-layered governance system where high-level policy, mid-level data integrity, and low-level operational security are deeply interconnected. For instance, a node's position in the ESPD (via its B and R scores) directly influences its Karma score (K_i), which in turn dictates its tolerance to security actions (its T_i). Simultaneously, a low D_t score degrades the node's reported B and Karma scores and signals a potential security incident, prompting a response from the KTSF system. This tight coupling ensures that the entire system remains coherent and responsive.

Despite the comprehensive nature of the proposed models, several areas require further clarification and refinement to move from theory to implementation. Key among these are the specific parameterization of the models. The weighting factors (w_V, w_m, \dots) in the Risk equation of the ESPD, the normalization schemes for translating raw ecotoxicity data into the $R_{\text{materials}}$ index, and the precise functional form of $f(K_i, B_i, R_i)$ for the tolerance radius T_i are not specified and will require careful calibration. Additionally, the exact algorithms for computing the diagnostic terms ($\Delta_{\text{drift}}, \Delta_{\text{var}}, \dots$) for the D_t scalar need to be defined with precision. Addressing these gaps is critical for the models to function as intended.

Based on this analysis, the following strategic recommendations are proposed for the development and implementation of the mathematical models:

1. Prioritize the Development of the

PlanningSafetySecurityAirWater2026v1.csv Shard Schema: The first concrete step should be to finalize the data schema outlined in this report. Creating a detailed specification for this shard forces clarity on data types, units, and dependencies, serving as the contract between all system components. This should be done in parallel with the development of the core C++ libraries.

2. Focus Initial Engineering Effort on the ESPD for the Phoenix Pilot: Given the immediate deadline for the Phoenix pilot, development resources should be concentrated on implementing and validating the ESPD model first. This involves coding the B and R calculations and using available Phoenix-specific data—particularly the grid intensity constraint of $\leq 50 \text{ g CO}_2/\text{kWh}$ —to define the initial "Deployable" zone. This will deliver the most urgent value.

- 3. Implement the D_t Module as a Real-Time, Parallel Service:** The Multonry Sensor-Trust Displacement logic should be developed as a separate, lightweight C++ library or service that runs concurrently with the main air-globe control loop. It should be designed to ingest raw sensor data, compute the trust scalar, and publish it to the shared `qpudatashard` for consumption by both the reporting and security modules. This modular approach enhances system resilience.
- 4. Define and Publish the "Eco-Grammar" for Fusion Logic:** The mathematical formulas and normalization schemes for combining disparate metrics (e.g., the weighted sum for risk, the convex fusion for Karma) constitute the system's "eco-grammar." This logic must be formally defined, documented, and published. This transparency is essential for third-party verification, auditing, and the future adaptation of the models by other regions.
- 5. Establish a Feedback Loop for Continuous Parameter Tuning:** The system must be designed to collect and report on the outcomes of the models (e.g., the number of nodes transitioning between ESPD zones, the frequency and cause of low- D_t events). This data should feed into a dashboard for operators, who can then iteratively tune the model parameters (weights, thresholds) to better match observed performance and desired policy goals.

By following these recommendations, the project can systematically transition from the high-level mathematical concepts presented here to a fully implemented, robust, and scalable system for the planning, safety, and security of cyboquatic air-globes, paving the way for their responsible and effective deployment.

Reference

1. ISO 14851:2019(en), Determination of the ultimate aerobic ... <https://www.iso.org/obp/ui/#iso:std:iso:14851:en>
2. biodegradation of synthetic polymers in <https://core.ac.uk/download/pdf/46919160.pdf>
3. (PDF) Critical Factors in Lab-Scale Compostability Testing https://www.researchgate.net/publication/384045396_Critical_Factors_in_Lab-Scale_Compostability_Testing

4. Transitioning to Microplastic-Free Seed Coatings <https://www.mdpi.com/2073-4360/16/14/1969>
5. Ics - Uci.edu DMDB Chandra Enron2.1 Words | PDF <https://www.scribd.com/document/675329216/Ics-uci-Edu-Dmdb-Chandra-Enron2-1-Words>
6. Test No. 201: Freshwater Alga and Cyanobacteria, Growth ... https://www.oecd.org/content/dam/oecd/en/publications/reports/2011/07/test-no-201-freshwater-alga-and-cyanobacteria-growth-inhibition-test_g1gh28f1/9789264069923-en.pdf
7. Toward the future of OECD/ISO biodegradability testing- ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC10033483/>
8. ISO 14851:2019 <https://www.iso.org/standard/70026.html>
9. Words | PDF | Science | Engineering <https://www.scribd.com/doc/88199361/Words>
10. THÈSE https://hal.science/tel-03121283v2/file/clement_jean-baptiste_these_version_finale.pdf
11. Actively deforming porous media in an incompressible fluid <https://arxiv.org/pdf/2107.03661>
12. Simulation-based high-speed elongational rheometer for ... <https://link.springer.com/article/10.1186/s13362-025-00168-x>
13. Hydro Economic Asymmetries and Common – Pool Overdraft in ... <https://agupubs.onlinelibrary.wiley.com/doi/full/10.1029/2022WR032136>
14. A conservative diffuse-interface method for compressible ... https://www.researchgate.net/publication/341762318_A_conservative_diffuse-interface_method_for_compressible_two-phase_flows
15. Statistical inference with a manifold-constrained RNA ... <https://www.nature.com/articles/s41592-024-02471-8>
16. chemtrain-deploy: A Parallel and Scalable Framework for ... <https://pubs.acs.org/doi/10.1021/acs.jctc.5c00996>
17. 机器学习2024_11_12 <https://arxivdaily.com/thread/61256>
18. Plug & play directed evolution of proteins with gradient- ... <https://iopscience.iop.org/article/10.1088/2632-2153/accacd>
19. Multi-Source Data Integration and Model Coupling for ... <https://www.mdpi.com/2227-9717/13/9/2998>
20. Computer Science <https://arxiv.org/list/cs/new>
21. ISO 14851:1999(en), Determination of the ultimate aerobic ... <https://www.iso.org/obp/ui/#iso:std:iso:14851:ed-1:en>
22. An electromagnetic strategy to improve the performance of ... <https://www.researchgate.net/publication/>

345389843_An_electromagnetic_strategy_to_improve_the_performance_of_PV_panel_under_partial_shading

23. Optimization of Chemical Processes <https://link.springer.com/content/pdf/10.1007/978-3-031-57270-8.pdf>
24. Algorithm-based fault tolerance applied to high ... https://www.researchgate.net/publication/223454627_Algorithm-based_fault_tolerance_applied_to_high_performance_computing
25. Advances in Atmospheric Sciences <https://iap.cas.cn/gb/xwdt/kyjz/202207/P020220711558871094364.pdf>
26. Reduced basis methods for urban air quality modeling <https://theses.hal.science/tel-01759410v1/file/TH2017PESC1230.pdf>
27. Modeling, optimization, and thermal management ... <https://www.sciencedirect.com/science/article/pii/S2590123025019954>
28. Mathematical Modelling of Environmental and Life ... [https://www.researchgate.net/profile/Ulrich-Ruede/publication/236892145_Analysis_of_a_Preconditioned\(CG\)_Method_for_an_Inverse_Bioelectric_Field_Problem/links/0a85e53664c38ba14b000000/Analysis-of-a-Preconditioned-CG-Method-for-an-Inverse-Bioelectric-Field-Problem.pdf](https://www.researchgate.net/profile/Ulrich-Ruede/publication/236892145_Analysis_of_a_Preconditioned(CG)_Method_for_an_Inverse_Bioelectric_Field_Problem/links/0a85e53664c38ba14b000000/Analysis-of-a-Preconditioned-CG-Method-for-an-Inverse-Bioelectric-Field-Problem.pdf)
29. Recovered_Sheet1 <http://www.lib.fzu.edu.cn/UploadFile/2019113152050211.xlsx>
30. Modeling security in cyber–physical systems <https://www.sciencedirect.com/science/article/abs/pii/S1874548212000443>
31. Developing Models for Physical Attacks in Cyber- ... https://www.researchgate.net/publication/320745177_Developing_Models_for_Physical_Attacks_in_Cyber-Physical_Systems
32. Formal Analysis of Security Properties of Cyber-Physical ... <https://www.semanticscholar.org/paper/Formal-Analysis-of-Security-Properties-of-System-on-Wang-Su/40b50054d4f87caed67f0b9d224003030256998c>
33. How to Deal with the Complexity of Future Cyber-Physical ... https://www.researchgate.net/publication/328467235_How_to_Deal_with_the_Complexity_of_Future_Cyber-Physical_Systems
34. Glider: Global and Local Instruction-Driven Expert Router <https://aclanthology.org/2025.emnlp-main.319.pdf>
35. A Survey of Vibe Coding with Large Language Models https://www.researchgate.net/publication/396498909_A_Survey_of_Vibe_Coding_with_Large_Language_Models/fulltext/68ef0fc4220a341aa1553778/A-Survey-of-Vibe-Coding-with-Large-Language-Models.pdf?origin=scientificContributions
36. Network Dictionary | PDF | Business | Computers <https://www.scribd.com/doc/309631022/Network-Dictionary>

37. Refugee-Related Sites on the World Wide Web <https://www.jstor.org/stable/pdf/45053048.pdf>
38. Corporate Social Responsibility in Tanzania: Experience of ... https://www.researchgate.net/publication/277385691_Corporate_Social_Responsibility_in_Tanzania_Experience_of_Misconception_Misuse_and_Malpractices_in_Rehema_Kilonzo_and_Tiina_Kontinen_eds_Contemporary_concerns_in_Development_Studies_Perspectives_form_
39. Beyond the Imitation Game Benchmark <https://arxiv.org/pdf/2206.04615>
40. Small-scale Freshwater Toxicity Investigations - Volume 2 <https://www.jlakes.org/ch/web/1-4020-3553-5.pdf>
41. Systematic Review of Toxicity Removal by Advanced ... <https://pubs.acs.org/doi/10.1021/acs.est.9b00570>
42. Adsorption + Theory, + Modeling, + And + Analysis. ... <https://www.researchgate.net/profile/Salah-N-Farhan/post/Does-the-adsorption-of-Cochineal-dye-onto-wool-follow-Pseudo-first-order-rate-kinetics/attachment/59d6374f79197b8077994ac0/AS%3A392500838715397%401470590900653/download/Adsorption+Theory%2C+Modeling%2C+And+Analysis.pdf>
43. Molecular simulation of fluid adsorption in flexible ... <https://theses.hal.science/tel-02181661/file/th%C3%A8se%20finale.pdf>
44. Big-Data Science in Porous Materials - ACS Publications <https://pubs.acs.org/doi/10.1021/acs.chemrev.0c00004>
45. Evaluation and Calibration of Coupled Low-Cost Particulate ... https://openscholarship.wustl.edu/cgi/viewcontent.cgi?article=2384&context=eng_etds
46. Freshwater Alga and Cyanobacteria, Growth Inhibition Test <https://www.oecd.org/content/dam/oecd/en/events/public-consultations/2025/12/first-commenting-round-for-the-draft-revised-test-guideline-201--freshwater-alga-and-cyanobacteria,-growth-inhibition-test/draft-revised-test-guideline-201-freshwater-alga-and-cyanobacteria-growth-inhibition-test.pdf>
47. Freshwater Alga and Cyanobacteria, Growth Inhibition Test https://www.oecd.org/en/publications/test-no-201-alga-growth-inhibition-test_9789264069923-en.html
48. A life cycle carbon dioxide equivalent emissions ... <https://www.sciencedirect.com/science/article/pii/S2590123024008442>
49. Estimating Carbon Dioxide Emissions from Power Plant ... <https://www.mdpi.com/2072-4292/16/7/1290>
50. Renewable power generation costs in 2024 https://www.irena.org/-/media/Files/IRENA/Agency/Publication/2025/Jul/IRENA_TEC_RPGC_in_2024_2025.pdf
51. THIRD NATIONAL COMMUNICATION OF SIERRA LEONE ... <https://unfccc.int/sites/default/files/resource/FinalThird%20Nat.%20Com.%20document%20111.pdf>

52. Processes, Volume 7, Issue 9 (September 2019) – 91 articles <https://www.mdpi.com/2227-9717/7/9>
53. China https://www.nsfc.gov.cn/Portals/0/fj/fj20240205_01.xlsx
54. Dynamic calibration of low-cost PM2.5 sensors using trust- ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC12228579/>
55. 3617623 12|3594859 time|3570693 year|3516017 publisher http://svn.apache.org/repos/asf/lucene/dev/tags/realtime_DWPT_final_2011-05-02/solr/src/test-files/Top50KWiki.utf8
56. cmnt_vocab.txt https://www.cs.cmu.edu/~ark/blog-data/data/blog_data_v1_0/dk/hbc_data/data/cmnt_vocab.txt
57. MONITORING, CONTROL AND EFFECTS OF AIR ... <https://pdfs.semanticscholar.org/4be8/1c30c8cc3d4e131ea076f44fb25638a86a11.pdf>
58. Labels and Definitions <https://www.eionet.europa.eu/gemet/exports/latest/gemet-definitions.html>
59. Issue 2 - Volume 100 - Physica Scripta <https://iopscience.iop.org/issue/1402-4896/100/2>
60. Evaluation of Speciation Technology https://www.oecd.org/content/dam/oecd/en/publications/reports/2001/04/evaluation-of-speciation-technology_g1gh1d17/9789264192782-en.pdf
61. (PDF) DPSIR Model Applied to the Remediation of ... https://www.researchgate.net/publication/343162158_DPSIR_Model_Applied_to_the_Remediation_of_Contaminated_Sites_A_Case_Study_Mar_Piccolo_of_Taranto
62. Nuclear fuel cycle and reactor strategies https://www-pub.iaea.org/MTCD/Publications/PDF/te_990_prn.pdf
63. Guidance on how aged sorption studies for pesticides should ... https://food.ec.europa.eu/system/files/2022-02/pesticides_ppp_app-proc_guide_fate_aged-sorption_0.pdf
64. and time-dependent sorption and desorption behavior of ... <https://pubmed.ncbi.nlm.nih.gov/20381112/>
65. Adsorption Time-Dependent changes in ... https://www.researchgate.net/figure/Adsorption-Time-Dependent-changes-in-adsorption-capacity-for-various-initial-MB_fig4_374715536
66. Test Guideline on the dustiness <https://www.oecd.org/content/dam/oecd/en/topics/policy-sub-issues/testing-of-chemicals/draft-test-guideline-dustiness-determination-of-manufactured-nanomaterials.pdf>
67. 2nd UNESCO Training Course on Sustainable ... <https://unesdoc.unesco.org/ark:/48223/pf0000153650>

68. overview-comments-received-guideline-environmental-risk ... https://www.ema.europa.eu/en/documents/comments/overview-comments-received-guideline-environmental-risk-assessment-medicinal-products-human-use-emea-chmp-swp-4447-00-rev-1_en.pdf
69. [2503.04954] Security-Aware Sensor Fusion with MATE <https://arxiv.org/abs/2503.04954>
70. (PDF) Trust-based security models in wireless sensor ... https://www.researchgate.net/publication/322150121_Trust-based_security_models_in_wireless_sensor_networks_a_survey
71. An Attack-Resistant Trust Model Based on ... https://www.researchgate.net/publication/273350314_An_Attack-Resistant_Trust_Model_Based_on_Multidimensional_Trust_Metrics_in_Underwater_Acoustic_Sensor_Network
72. An Efficient Distributed Trust Model for Wireless Sensor ... https://www.researchgate.net/publication/270793788_An_Efficient_Distributed_Trust_Model_for_Wireless_Sensor_Networks
73. Artificial Intelligence in Healthcare: 2020 Year in Review https://www.researchgate.net/profile/Piyush-Mathur-2/publication/349570341_Artificial_Intelligence_in_Healthcare_2020_Year_in_Review/links/603676e2a6fdcc37a84d874c/Artificial-Intelligence-in-Healthcare-2020-Year-in-Review.pdf
74. 333333 23135851162 the 13151942776 of 12997637966 <ftp://ftp.cs.princeton.edu/pub/cs226/autocomplete/words-333333.txt>
75. Depletion of Mettl3 in cholinergic neurons causes adult- ... [https://www.cell.com/cell-reports/fulltext/S2211-1247\(24\)00327-9](https://www.cell.com/cell-reports/fulltext/S2211-1247(24)00327-9)
76. Identification of high-confidence human poly(A) RNA isoform ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC8906549/>
77. Reading canonical and modified nucleobases in 16S ... <https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0216709>
78. Review RNA vaccines for cancer: Principles to practice <https://www.sciencedirect.com/science/article/pii/S1535610824001685>
79. CPVM-extending PVM for consistent checkpointing <https://ieeexplore.ieee.org/document/500570/>
80. CPVM - Extending PVM for Consistent Checkpointing <https://ieeexplore.ieee.org/iel3/3578/10690/00500570.pdf>